


# Fuzzing, DevOps, and You: Getting to Know Fuzzing

By  
Allison Marie Naaktgeboren  
[naak@pdx.edu](mailto:naak@pdx.edu)





# About Allison

- There's more than one Allison Naaktgeboren!
  - Publication Name: Allison Naaktgeboren
- New Graduate Research Assistant, Portland State University
  - Security Lab, advised by Dr. Wu-chang Feng
  - Studying influence of the initial seed corpus on fuzzer performance
- Old Code Monkey
  - Signal Sciences, Mozilla, FactSet, Amazon, Cisco, RI Biorobotics Laboratory, Coding with Kids
  - Carnegie Mellon University, SCS, BS in CS
    - Research in robotics
- Community: OWASP PDX, WWC PDX, First Robotics, Blackhoodie
- Speaking: DEF CON HHV, BSides PDX, Linux Security Summit, Shakticon
- CTF
  - President, cofounder PSU Security Club, void\* vikings
  - Captain, cofounder QultoftheQuantumQapybaras
  - (Useless) Minion, Samurai



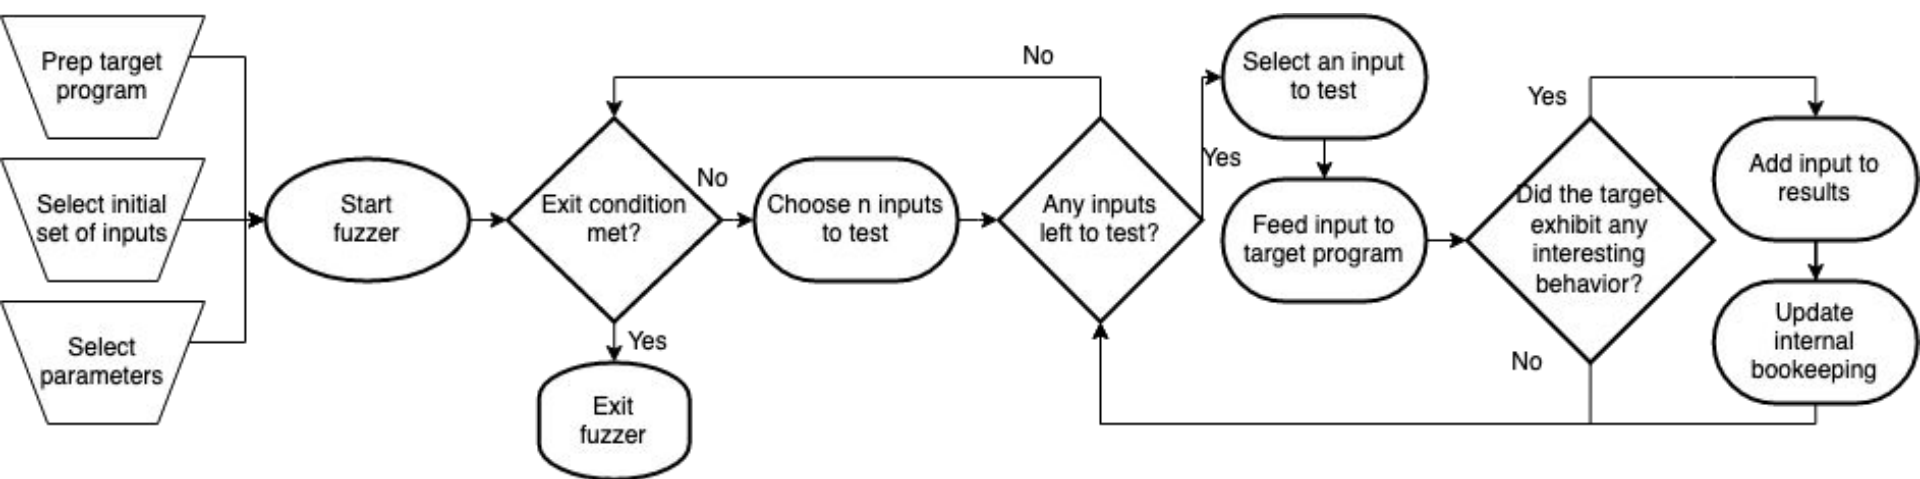
*DC29 Finals, credit: px*

What's this new fuzzing  
thing I keep hearing about  
???



# Fuzzing, Fuzz Testing

- A dynamic stochastic software-testing technique
- Its goal is to thoroughly explore the input space of the target program looking for inputs (seeds) that cause *interesting* behavior
  - 'Interesting' is not universally defined
  - Symbolic/concolic execution(e.g. angr) is sometimes considered part of fuzzing
- First published by Bart Miller & students in December, 1990
  - [An empirical study of the reliability of UNIX utilities | Communications of the ACM](#)
  - Magazine article
- Terminology Cheat sheet:
  - <https://github.com/anaaktge/talks/blob/main/fuzzingtalkglossary.md>




- ❑ Setup
  - ❑ Prep target program
  - ❑ Select initial inputs (seeds)
  - ❑ Select parameters
- ❑ Start fuzzer
- ❑ While exit condition is not satisfied:
  - ❑ Get a set of inputs (seeds) to test
  - ❑ For each input in set:
    - ❑ Feed target program input and run
    - ❑ If interesting behavior happens
      - ❑ Store input and result
      - ❑ Update bookkeeping
    - ❑ Else move on

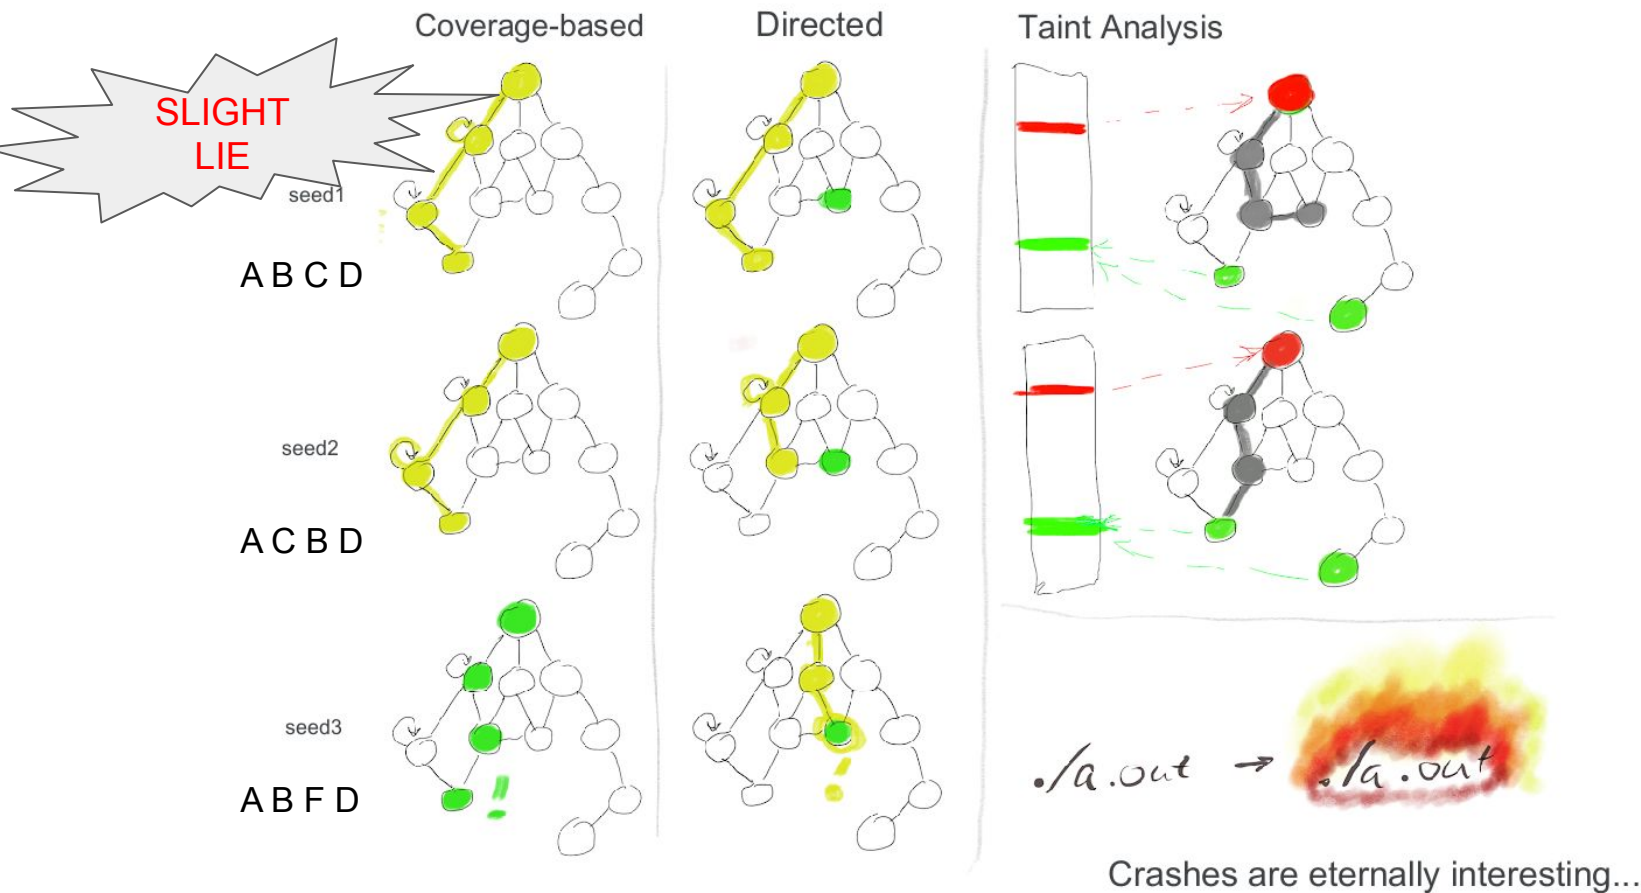
# Understanding Fuzzer Differences

1. How do you define interesting?
2. How do you get your seeds (inputs)?
3. What's your exit condition?
4. How much do you know about the target source?
5. What's your focus/speciality?
6. Do you have a secret sauce?

# 1 - How do you define interesting?



# 1 - How do you define interesting?





## 2 - What's your exit condition?



Scheduled Duration



Out of inputs to test




User Initated Exit



One and Done

### 3 - Where do you get your seeds (inputs) ?

Starting  
Seed




Mutation

Starting  
Algorithm



Generation

Starting  
???



Classic, Naive

# 4 - What do you know about the target source?

```
105 // returns a reference
106 fn get_info(&self, id: usize) -> Result<String, String> {
107     dbg!("get info on id: {}", id);
108     if id < self.functions.len() {
109         let f = &self.functions[id];
110         let basic_info = f.to_string();
111         let callees = f.calls_ids().fold("Makes calls to: ".to_string(), |a, &i| {
112             a + "(fn " + &i.to_string() + ": " + self.functions[i].to_string() + ") ";
113         });
114         let body = if let Some(b) = f.body_as_ref() {
115             b.to_string()
116         } else {
117             "[no body; this is an import]".to_string()
118         };
119     }
```

00003780:	2434	6472	6f70	3137	6836	3365	6234	3832
00003790:	3437	3766	3433	6166	3345	005f	5a4e	3838
000037a0:	5f24	4c54	2468	6173	6862	726f	776e	2e2e
000037b0:	7363	6f70	6567	7561	7264	2e2e	5363	6f70
000037c0:	6547	7561	7264	244c	5424	5424	4324	4624
000037d0:	4754	2424	7532	3024	6173	2475	3230	2463
000037e0:	6f72	652e	2e6f	7073	2e2e	6472	6f70	2e2e
000037f0:	4472	6f70	2447	5424	3464	726f	7031	3768
00003800:	6139	6163	3730	6433	3663	3164	3038	6331
00003810:	4500	5f5a	4e39	305f	244c	5424	6861	7368
00003820:	6272	6f77	6e2e	2e73	636f	7065	6775	6172
00003830:	642e	2e53	636f	7065	4775	6172	6424	4c54
00003840:	2454	2443	2446	2447	5424	2475	3230	2461

Whitebox

00003780:	2434	6472	6f70	3137	6836	3365	6234	3832
00003790:	3437	3766	3433	6166	3345	005f	5a4e	3838
000037a0:	5f24	4c54	2468	6173	6862	726f	776e	2e2e
000037b0:	7363	6f70	6567	7561	7264	2e2e	5363	6f70
000037c0:	6547	7561	7264	244c	5424	5424	4324	4624
000037d0:	4754	2424	7532	3024	6173	2475	3230	2463
000037e0:	6f72	652e	2e6f	7073	2e2e	6472	6f70	2e2e
000037f0:	4472	6f70	2447	5424	3464	726f	7031	3768
00003800:	6139	6163	3730	6433	3663	3164	3038	6331
00003810:	4500	5f5a	4e39	305f	244c	5424	6861	7368
00003820:	6272	6f77	6e2e	2e73	636f	7065	6775	6172
00003830:	642e	2e53	636f	7065	4775	6172	6424	4c54
00003840:	2454	2443	2446	2447	5424	2475	3230	2461
00003850:	7324	7532	3024	636f	7265	2e2e	6f70	732e
00003860:	2e64	6572	6566	2e2e	4465	7265	6624	4754
00003870:	2435	6465	7265	6631	3768	3931	6636	6163
00003880:	6537	3233	3864	6331	6231	4500	5f5a	4e39
00003890:	305f	244c	5424	6861	7368	6272	6f77	6e2e
000038a0:	2e73	636f	7065	6775	6172	642e	2e53	636f
000038b0:	7065	4775	6172	6424	4c54	2454	2443	2446


Blackbox

Graybox


5 - What's your focus or speciality?



## 6 - Do you have a secret sauce?



## REDQUEEN: Fuzzing with Input-to-State Correspondence



## FAIRFUZZ: A Targeted Mutation Strategy for Increasing Greybox Fuzz Testing Coverage

# Understanding Fuzzer Differences tl;dr

- How do you define interesting?
  - Coverage based
  - Directed
  - Taint Analysis
  - Crashes
- What's your exit condition?
  - Time out
  - ^C
  - Out of stuff to test
  - first crash found
- How do you get your seeds (inputs)?
  - Mutational
  - Generational
  - Classic (naive, random)
- How much do you know about the target source?
  - White box - full knowledge, source
  - Greybox - instrumented binary
  - Blackbox - no knowledge
- What's your focus/speciality?
  - Protocols, Kernels
  - Webpages, IoT
  - cars
- Do you have a secret sauce?
  - Super Awesome Optimization™
  - Combined with symbolic execution (hybrid fuzzing)
  - ML/AI/DNN



# Who's at the Fuzzing Party?

...and why are some of them really angry?



# Fuzzing Party: Academic Researchers

You might be a...Phd Student, Faculty, Masters Student

Your goals might be...

- A previous group's goals
- Studying Software Testing
- Studying another aspect of CS

Features to look for...

- Benchmark
- Good substrate for testing new ideas
- Publishable
- Citable

Fuzzers you might like...

- [AFL](#)
- Developing your own to advance the State of the Art (SOTA)



# Fuzzing Party: Offensive Security

You might be a... Vulnerability Researcher, Penetration Tester, Bug Bounty Hunter

## Your goals are...

- Finding at least one vulnerability (exploitable bug) or chain
- ASAP!
- Preferably before anyone else
- Don't care about 'why'

## Features to look for...

- Vulnerabilities
- Fast to first find
- Find tricky bugs other fuzzers have missed
- Source code not required

## Fuzzers you might like...

- [libFuzzer](#)
- [MOpt-AFL](#)

# Fuzzing Party: Defensive Security

You might be a... Blue teamer, Application Security Engineer, Security Consultant

Your goals are...

- Find all the vulnerabilities
- Find them all before release
- Care about 'why'

Features to look for...

- Probe as much of the product as possible
- Complement other security practices, tools

Fuzzers you might like...

- [AFLplusplus](#)
- [honggfuzz](#)
- [AFLGo: Directed Greybox Fuzzing](#)

# Fuzzing Party: Software Owners

You might be a... Engineering Manager, Lead Software Engineer, QA Engineer

## Your goals are...

- Find vulnerabilities
- Find all the bugs
- Find them before release
- Really care about 'why'

## Features to look for....


- Target specific commits
- CI/CD integration
- Stability
- Finds more than just vulnerabilities
- Code coverage might interest you

## Fuzzers you might like...

- [classic fuzz](#)
- [TOFU:](#)  
[Target-Oriented](#)  
[FUzzer](#)
- Microsoft OneFuzz

# Fuzzing Limitations

- Hard on your machine
  - CPU intensive, RAM intensive
  - Consumes a lot of power and \$\$
- Relies on probability.
  - Probability is not always your friend
- Fuzzing does not tell you the 'why'
  - Understanding the results can be nontrivial
- Subproblems in fuzzing are nontrivial
  - Code coverage algorithms are not absolute
- Figuring out how long to run a fuzzer is nontrivial
  - When does it stop making progress?
  - Going minimum is 24 hours
  - Magma results suggest that a week (168 hours) might be better



*Sometimes, you're just not a good fit..*

# Is it a Good Fuzzing Target?

## Yes

- Obvious parsers
- Stealth parsers (load/store)
- Command line utilities
- Text based APIs
- Libraries
  - images

## No

- GUI only access
  - Files are fine
- Difficult for probability
  - Deep, complex, non visible state
  - Complex sequences & exchanges
- Classic stumbling blocks
  - Kernels
  - Network protocols
  - Magic #s, Tight branches

# DevOps Specific Advice

- VMs will need to be beefy in **memory AND CPU**





- Allow a lot of memory consumption
  - Probably going to want to get as close to bare metal as you can
  - Don't start with distributed fuzzing
- May need more than one fuzzer
  - Start small, high value, easy fit
- Definitely need more than one harness and corpus
  - You've already got the corpus
- May not fit directly into a CI/CD pipeline (out of band testing)
- May want to add in auto-testing of crashes

# Acknowledgements & Further Reading

- Dr. Wu-chang Feng, advisor
- National Science Foundation
- The Security Lab, Portland State
- DevSLOP organizers

- Invention of fuzzing: [An empirical study of the reliability of UNIX utilities | Communications of the ACM](#)
- 30 years later: [The Relevance of Classic Fuzz Testing: Have We Solved This One?](#)
- [Magma Fuzzing Benchmark](#)
- FairFuzz: [FairFuzz: a targeted mutation strategy for increasing greybox fuzz testing coverage](#)
- AFLFast: [AFLFast \(extends AFL with Power Schedules\)](#)
- [REDQUEEN: Fuzzing with Input-to-State Correspondence](#)





## Bonus Slides



# AFL (American Fuzzy Lop)

- Famous in research and industry
  - <https://lcamtuf.coredump.cx/afl/>
- Characteristics
  - Target Knowledge: Greybox
  - Seed Generation: Mutational
  - Interesting: Code coverage
  - Interesting: Crashes, hangs
  - Exit condition: time, ^C, empty
  - Driver: AFL speaks only stdin(ish)
- Lots of knobs to tune
  - Use them!
  - **An hour to start, a lifetime to master**
- Posse of helper tools
  - afl-min, afl-plot, etc


american fuzzy lop 0.47b (readpng)			
<b>process timing</b>		<b>overall results</b>	
run time : 0 days, 0 hrs, 4 min, 43 sec		cycles done : 0	
last new path : 0 days, 0 hrs, 0 min, 26 sec		total paths : 195	
last uniq crash : none seen yet		uniq crashes : 0	
last uniq hang : 0 days, 0 hrs, 1 min, 51 sec		uniq hangs : 1	
<b>cycle progress</b>		<b>map coverage</b>	
now processing : 38 (19.49%)		map density : 1217 (7.43%)	
paths timed out : 0 (0.00%)		count coverage : 2.55 bits/tuple	
<b>stage progress</b>		<b>findings in depth</b>	
now trying : interest 32/8		favored paths : 128 (65.64%)	
stage execs : 0/9990 (0.00%)		new edges on : 85 (43.59%)	
total execs : 654k		total crashes : 0 (0 unique)	
exec speed : 2306/sec		total hangs : 1 (1 unique)	
<b>fuzzing strategy yields</b>		<b>path geometry</b>	
bit flips : 88/14.4k, 6/14.4k, 6/14.4k		levels : 3	
byte flips : 0/1804, 0/1786, 1/1750		pending : 178	
arithmetics : 31/126k, 3/45.6k, 1/17.8k		pend fav : 114	
known ints : 1/15.8k, 4/65.8k, 6/78.2k		imported : 0	
havoc : 34/254k, 0/0		variable : 0	
trim : 2876 B/931 (61.45% gain)		latent : 0	

The iconic main screen

Credit: <https://lcamtuf.coredump.cx/afl/>

# AFL Gotchas

- Start
  - Don't have source to instrument
- Prep
  - Didn't create a good initial corpus
    - One call is unlikely to probe the entire program
  - Didn't instrument everything in binary (including the shared libraries)
- AFL Parameter Configuration
  - Didn't tune parameters
  - Abused the @ parameter
- Running AFL
  - Need root access to set processor settings
  - Didn't run the fuzzer long enough (1 to 7 days)
  - Might need to repeat the fuzzer run
- Reviewing, Analyzing the Results
  - Crash != vulnerability.
    - Might not have security value
  - Unique crashes may not be unique
    - 20+ crashes might only be 1 source



# How do I get a good initial seed corpus for AFL ?

- ❖ We don't really know what an 'optimal' corpus is in the general case
- ❖ Get to know the target:
  - Run the target a few times
  - Examine source or binary
- ❖ Things I've done to get an initial corpus
  - Raid the unit test/test cases for the project
  - Run 'strings' on the target binary and edit it down
  - Run the program a few times and find a few inputs
  - Do some RE (Reverse Engineering) on the target!
- ❖ Running afl-min is a very good idea!



*wikiMedia, navy bean*