


PIPE Cleaner:

Configurable PIPE-Based Policy-Informed Root-Cause Fuzzing

Allison Naaktgeboren
Portland State University



"You must be this tall to find this bug"

Photo: Allison

Terminology Cheat sheet:

<https://github.com/anaaktge/talks/blob/main/fuzzingtalkglossary.md>

Agenda



Goal: Help you understand...

1. Fuzzing holistically
2. Where PL might help
3. Maybe my own research

Topics

- About Me
- Motivation & What is fuzzing ?
- What is the root cause problem and who cares?
- What is PIPE & TaggedC?
- POC: PVI Violation
- PIPE Cleaner: The Fuzzer I hope to build

Not a [tenure-track] Professor!

Currently: Draper Labs Intern, working on VMF fuzzer

Doctoral Researcher, Portland State, 2020-Present (MS, cybersec)

Under Dr. Andrew Tolmach (QuickChick, SF coauthor)

RPE (Quals) “Sowing the Seeds of Fuzz: Does the Influence of the Initial Seed Corpus Follow a Universal Law?” Jun 22

CTF teams: void * vikings, founder + advisor. 侍 minion

Instructor, Lecturer, Intro to Systems (+ all legal obligations)

BS in CS Carnegie Mellon University, 2004-2008 (SCS '08)

Undergraduate research in robotics, roboclub officer

Founding TA 15-440, 15-213, 15-111, 16-311

(Senior) Code Monkey, Hiring Manager, etc ~2009-2019

Cisco, Amazon, Factset, Mozilla (Firefox), Signal Sciences (Fastly)



*The best regalia accessories
are minions who bring you
donuts and coffee
Photo: kbrosnan*

My prior PL background...



Motivation: Vulnerabilities Suck (and cost \$\$\$)

Security Bulletin: IBM Security Guardium is affected by Open Source libxml2 vulnerabilities

Amazon Linux Security Center

Security Bulletin: Vulnerability in libxml2 affects IBM InfoSphere Streams. (CVE-2015-8317)

Heartbleed bug 'will cost millions'



From my RPE and guest lecture

How much will it cost to secure open-source software? OpenSSF says \$147.9M

Affected Packages

Platform
Amazon Linux 2
Amazon Linux 1

Automation to the Rescue? Fuzzing Better than Santa Claus?



Security

Linus Torvalds lauds fuzzing for improving Linux security

But he's not at all keen on Santa Claus or fairies

By [Simon Sharwood](#), APAC Editor 16 Oct 2017 at 07:03

ZDNet



MENU



US

MUST READ A POPULAR VIRTUAL KEYBOARD APP LEAKS 31 MILLION USERS' PERSONAL DATA

Linux security: Google fuzzer finds ton of holes in kernel's USB subsystem

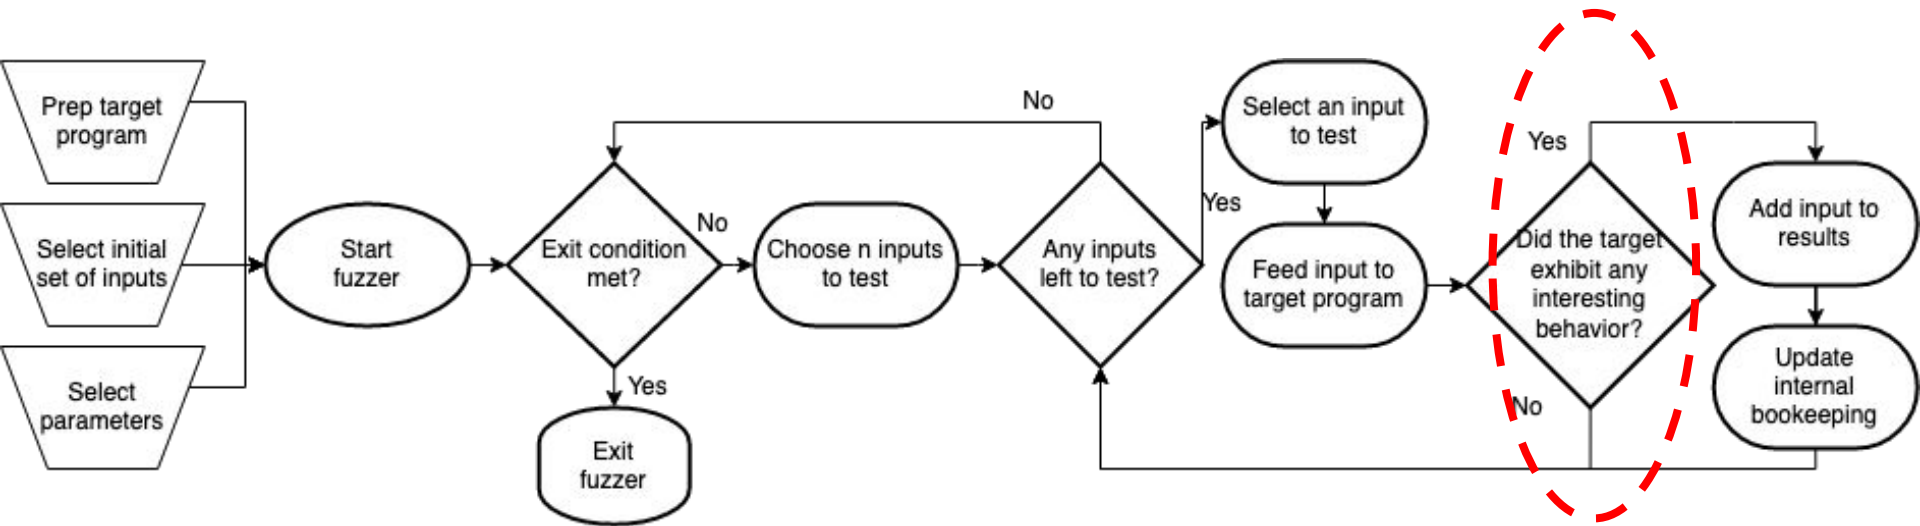
A Google-developed kernel fuzzer has helped locate dozens of Linux security flaws.



By [Liam Tung](#) | November 8, 2017 -- 12:43 GMT (04:43 PST) | Topic: [Security](#)

Overview of Fuzzing

Does not include why the target did something interesting!



- Aims to thoroughly explore the input space of the fuzzee (victim, target) looking for inputs (seeds) that cause interesting behavior
- Definition of “interesting” varies. Classic definition is crashes
- A stochastic (probabilistic) dynamic software-testing technique
- Property based testing, symbolic/concolic execution are arguably subclasses


The Problem of Root Cause

“Developers generally appreciate bug reports, but they can sometimes be a bit **less enthusiastic** about a **flood of reports from automated fuzzing systems**.” <https://lwn.net/Articles/904293/>

“..you have an ethical and moral responsibility to do some of the work to narrow down and **identify the cause of the failure**, not just throw them at someone to do all the work”
<https://lwn.net/Articles/917762/>

How bad is it? (audience participation)
If a fuzzer produced 3,020 crashes, how many bugs do you think there were?

**Ans: there were 15 bugs
(FuzzerAid)**



Not Everyone Cares about Root Cause

Offensive (Red):

Vulnerability Researcher,
Penetration Tester, Bug
Bounty Hunter

Goals:

- Finding at least one vulnerability or chain
- ASAP! Preferably before anyone else
- Rarity, stability good
- **Don't care about 'why'**

Defensive (Blue):

Application Security
Engineer, Security
Consultant

Goals:

- Find all the vulnerabilities
- Find them all before release
- **Care about 'why'**

Software Owners:

Engineering Manager,
Software Engineer,
QA Engineer

Goals:

- Find all vulnerabilities
- Find all bugs
- Find them before release
- **Really care about 'why'**

Academics: Phd

Student, Faculty,
Masters thesis

Goals:

- A previous group's goals
- Studying Software Testing
- Studying another aspect of CS
- Clear, easy metrics

PIPE (Programmable Interlocks for Policy Enforcement)

- Is a Tag based hardware security reference monitor
 - ISA extension
- A tag represents arbitrary metadata associated with some data
 - Ex “This is a pointer to object a”, “This is low security data”, “The active function is f”, “this instruction can only be used by high security data’
- All data have tags
 - Check relevant tag rules on each execution step
 - if policy is violated (maybe) failstop
- PIPE, doesn’t care about language, high or low
- Problem: hardware doesn’t entirely exist yet, and what does isn’t widely available .

THE



PROGRAMMING LANGUAGE

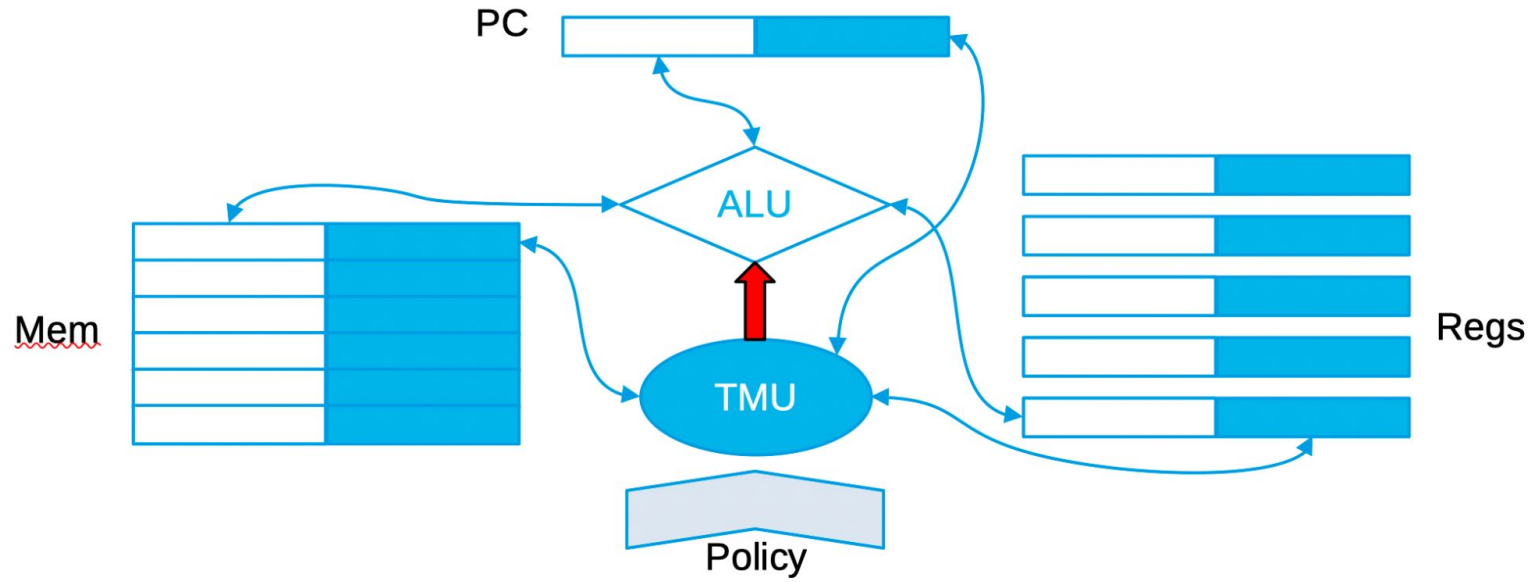
the leading cause of security analysts

TaggedC

- Built on Compcert's Interpreter
 - Give semantics for all undefined behavior
- PIPE policies are encoded at source level
 - Developers rarely speak assembly or machine code
- PL-y magic happens here
 - Sean Anderson's dissertation work
- Formalizing Stack Safety as a Security Property
 - <https://arxiv.org/abs/2105.00417> preprint
 - Sean is presenting it at a conference later this month
 - We have another paper in submission which discusses TaggedC in more detail
- Verified compilation part
 - Towards formally verified compilation of tag-based policy enforcement
 - <https://dl.acm.org/doi/10.1145/3437992.3439929>

Possible PIPE Implementation

Metadata Tagging in PIPE



TaggedC POC Policy Violation

- Provenance via Integer (PVI)
 - Supports arbitrary integer arithmetic on (int-cast) pointers
 - Does not support crossing object bounds
 - Memory model from optimization literature
- Why would you do this in C?
 - Use lower order bits as flag
 - E.g. Cheney's garbage collection algorithm

```
int main() {  
    int x[10];  
    int y[10];  
  
    *(x + 10) = 42;  
    return y[0];  
}
```

```
a@pyrite tagged-c % ./violation1  
zsh: abort ./violation1
```

VS

```
a@pyrite tagged-c % ./ccomp -p pvi violation1.c  
Failstop on policy PVI::LoadT X Failure
```

The Dissertation Fuzzer I Hope To Build: PIPE Cleaner

- Assuming I have cleverly designed policies in TaggedC
 - Specifying these formally is a nontrivial PL sort of challenge
 - especially beyond memory corruption
 - Formal speak: interested in Hyperproperties
- Fuzzer searches for PIPE Policy violations rather than segfaults
 - Fuzzer has root cause at time of fault, not something you get from a segfault
 - Encoding “interesting behavior” as policies means I can change what to fuzz for without changing the fuzzer
 - Policies can better represent things like SQL injection that don’t make sense as segfaults
- Ideally, first fuzzer to have no false positives by construction
 - Imagine instead of 3,020 crash reports, you only have to read 15 reports?
 - instead of consigning fuzzer report to oblivion or summer interns, bugs might get fixed
 - Fuzzing is always unsound (false negatives. May miss bugs)

Thank you for Listening <3



Questions & Unsolicited Advice

1. Lock/close your fsck-ing laptops!
 - a. You wash your hands after you use the bathroom, right?
 - b. It's a possible FERPA violation if you teach in the US
2. Use MFA, especially on school email
3. Your (risky) behavior matters more than OS.
 - a. Stay Patched.
4. **Vulnerabilities are expensive, your mistakes (or laziness) are free!**
5. Have reason for students to learn PL beyond “purity”.
I had to wait 18 years for a reason. They won't.